



Vite



The Sharding Solution for MySQL 8

Liz van Dijk - @lizztheblizz

Who am I?



Liz van Dijk - Solution Architect

- Howest University > Percona > PlanetScale
- New to Vitess, MySQL has been my world for 8+ years
- liz@planetscale.com
- @lizztheblizz



PlanetScale

- Founded in February 2018
- Venture backed: a16z, SignalFire
- 35 employees mostly in Mountain View, CA

It's 2020. I blinked, and...



- Containers everywhere. Stateless workloads reign supreme.
- Public Cloud diversification. GCP and Azure are fully fledged counterparts to AWS.
 - Competition is fierce and vendor lock-in is a real risk.
- Container provisioning and orchestration are just as common as full VM. Kubernetes appears to be the new standard.



What is Vitess?



Vitess

**Cloud
Native**

**Massively
Scalable**

**Highly
Available**

**Based on
MySQL**

Vitess Stats



Started
2010



CNCF
Graduated



200
Contributors



9,500
Stars

v5.0
(4 Feb 2020)



18,000
Commits

1100
Slack Members

1200
Forks



Key Adopters



* Employing Active project maintainers

Vitess' Vital, but Vexing Vocabulary




- Think of Vitess as an extension bolted on to MySQL, both from the Dev and the Ops points of view.
- Various components coming together
- We will break down Infrastructural and Logical concepts



Tablet




- Most basic “worker” unit of a Vitess Cluster
- MySQL Server may be any flavour
- VTTablet** is a *sidecar process*
- Tablets can fulfill multiple roles
 - Master
 - Failover Replica
 - Analytics Replica



Keyspace & Shards




- ❖ **Keyspace** is an analog to what we call a *logical database*.
- ❖ A single keypace can be logically divided into **Shards**, internally identified by a **Keyspace ID**.
- ❖ Keyspaces don't need to be sharded.
- ❖ A Keyspace Shard will contain 1 Master Tablet and as many Replica Tablets as desired.






- Infrastructural Failure Domain, think AZ's
- Examples:
 - a full data center
 - a subset of a data center
 - a Kubernetes cluster





- **VTGate** is a stateless proxy, and the entry point into the cluster
 - ◻ Speaks MySQL protocol as well as gRPC
 - ◻ Can be connected to and presents the cluster as a monolithic database
 - ◻ Interprets SQL and supports Vitess-specific hints
 - ◻ Considers the cluster topology and routes requests accordingly

Architecture



VSchema



- Extension of the database schema as we know it
- JSON format
- Describes sharding logic & additional Vitess functionality like sequences

```
{
  "sharded": true,
  "vindexes": {
    "hash": {
      "type": "hash"
    }
  },
  "tables": {
    "user": {
      "column_vindexes": [
        {
          "column": "user_id",
          "name": "hash"
        }
      ]
    }
  }
}
```

VIndex



- Builds the bridge between a table's column values and the correct Keyspace ID
- Sharded tables should have at least a Primary **VIndex**, but many secondary VIndex types are available.

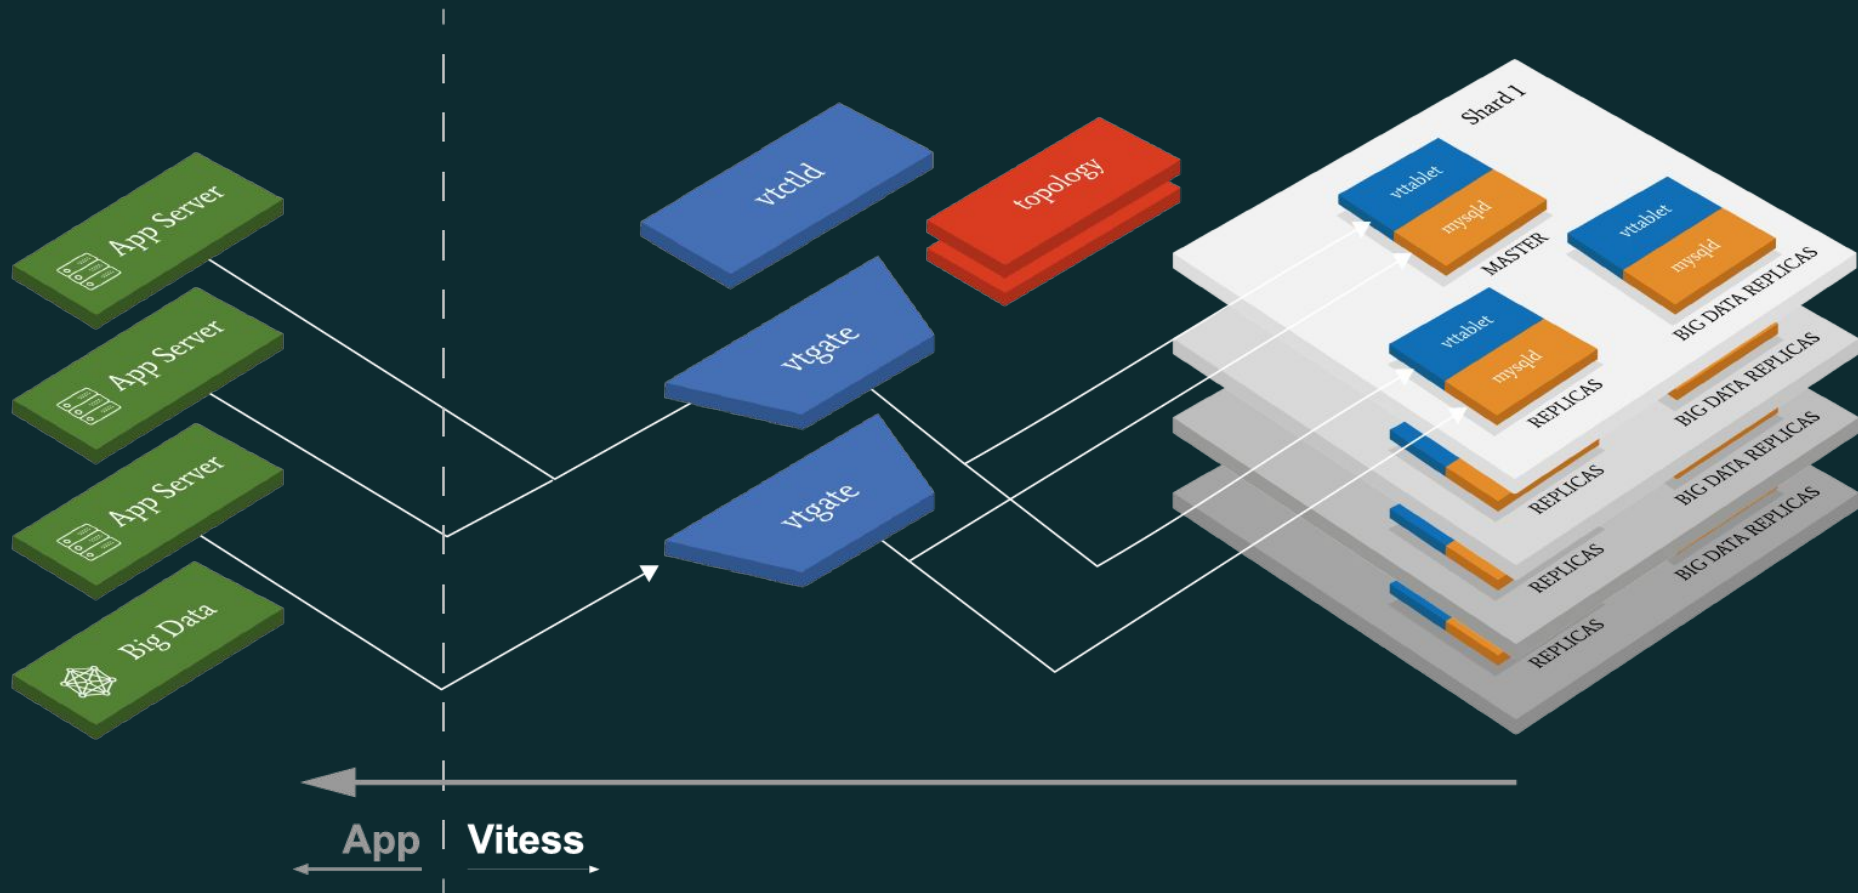
```
{
  "sharded": true,
  "vindexes": {
    "hash": {
      "type": "hash"
    }
  },
  "tables": {
    "user": {
      "column_vindexes": [
        {
          "column": "user_id",
          "name": "hash"
        }
      ]
    }
  }
}
```



@lizztheblizz

DEMO

Architecture



One more thing...



- Running Vitess on Kubernetes just got a whole lot easier
 - <https://github.com/planetScale/vitess-operator>
- We eat our own dog food and make sure it stays fresh.
 - PlanetScale CNDb runs on this operator
 - We have 3 dedicated contributors
- Eager to get feedback. Please try it out and tell us all about it!

Call us, maybe?



- ◆ Project website: vitess.io
- ◆ Official docs: vitess.io/docs
- ◆ Slack Community: vitess.io/slack - Highly recommended!

Come say hi at:

- ◆ SCALE18x
- ◆ KubeCon EU
- ◆ Percona Live Austin ... ?



@lizztheblizz

Questions?

Concepts: Sharding Functions



Or, add your own custom sharding function!

binary	Identity
binary_md5	md5 hash
hash	3DES null-key hash
numeric	Identity
numeric_static_map	A JSON file that maps input values to keyspace IDs
unicode_loose_md5	Case-insensitive (UCA level 1) md5 hash
reverse_bits	Bit Reversal